I2I Agentic AI Starter Guide

Complex Multiple-Choice Assessment (50 Questions)

Instructions: Choose the single best answer for each question. Answers and a quick key appear at the end of this document. Recommended pass bar: 40/50 (80%) before claiming readiness for design discussions.

Q1. A team states 'we use Agentic AI' in an architecture review but cannot describe autonomy, planning, or tool choice. What is the strongest critique?

A. Agentic AI is not a recognized term
B. Agentic capability is a spectrum; the claim lacks architectural specificity
C. Only multi-agent systems may use the term
D. LLMs cannot participate in agentic systems

Q2. Which statement best separates agent, workflow, and platform?

A. They are interchangeable synonyms
B. Agent decides/acts; workflow structures steps and state; platform helps build and run them
C. Platform always means the LLM vendor only
D. Workflow is only the UI theme

Q3. A multi-agent pipeline always executes Agent A then B then C with no runtime decision-making. How should this be classified?

A. Strongly agentic AI by definition
B. Primarily workflow automation, even though multiple agents exist
C. Pure A2A mesh
D. An MCP host

Q4. When is adding a second agent least justified?

A. Legal and finance need separate ownership and systems
B. A vendor agent must remain in another trust boundary
C. One agent with a few well-scoped tools already completes the task reliably
D. Specialists require independent release cycles

Q5. For an enterprise leave system that must authenticate, enforce blackout dates, submit to HRIS, then optionally draft a manager note, which pattern fits best?

A. Fully free planner with no fixed steps
B. Hybrid: deterministic rails for identity/policy/submit; LLM optional for drafting
C. Five debating A2A agents with no schemas
D. Memory-only design with no tools

Q6. What is the primary advantage of workflow-driven orchestration?

A. Maximum path surprise
B. Predictability, auditability, and testable branches
C. Eliminating all tools
D. Removing the need for AuthZ

Q7. Autonomous LLM orchestration is most appropriate when:

A. Every step is legally prescribed and identical
B. The task space is open-ended and useful paths cannot all be pre-encoded
C. Tools must never be called
D. Observability is intentionally disabled

Q8. An orchestrator's main responsibility is to:

A. Replace all tools with one megaprompt
B. Decompose goals, delegate, merge results, and decide stop or escalate

C. Store API keys in the system prompt
D. Retrain the base model weekly

Q9. Which design-review question most directly reduces production incidents?

A. Which emoji should the agent use?
B. What is in state after each step, and what happens on tool failure, retry, or HITL?
C. Can we rename the agent to sound more autonomous?
D. Which font does the chat UI use?

Q10. Why do tool descriptions and parameter schemas matter so much?

A. They are ignored by modern models
B. They heavily influence tool selection and argument correctness
C. One do-everything tool is always superior
D. Schemas replace authentication

Q11. A model requests delete_user with a guessed id. Which control is least sufficient alone?

A. Server-side AuthZ, allowlists, and HITL for destructive tools
B. Stronger wording in the system prompt only
C. Schema validation and trusted identity from session context
D. Audit logging of tool calls

Q12. In a ReAct-style loop, 'Observe' means:

A. Looking at the user only once at start
B. Incorporating tool or environment results into the next reasoning step
C. Retraining the foundation model
D. Deleting all memory

Q13. Why might a system use a small model for routing and a larger model for planning?

A. MCP requires dual models
B. Cost and latency tradeoffs: cheap classification vs stronger multi-step reasoning
C. Small models cannot call tools
D. Large models cannot follow schemas

Q14. Critical authorization for submit_leave should primarily live in:

A. Only the system prompt
B. The application/tool layer with real AuthN/AuthZ checks
C. The marketing site copy
D. The model card PDF alone

Q15. LLMs need external memory primarily because:

A. They automatically remember all users forever
B. They are effectively stateless across calls unless context is provided or stored
C. Memory replaces tools
D. Memory removes the need for guardrails

Q16. Storing sensitive medical leave reasons in long-term free-text memory without retention controls mainly risks:

A. CSS layout breakage
B. Privacy and compliance exposure
C. TCP window sizing
D. Font licensing

Q17. Which placement best prevents a disallowed tool call even if the model requests it?

A. Output style guide only
B. Tool gateway enforcing allowlist, AuthZ, and schema validation before execution
C. Increasing temperature
D. Longer system prompt with more adjectives

Q18. HITL is best described as:

A. Proof the agent failed
B. An intentional pause for human approval on high-impact actions
C. A replacement for all tools
D. Only suitable for demos

Q19. An agent may draft external emails freely but must get approval before sending. This illustrates:

A. No guardrails
B. Risk-tiered control: looser on reversible drafts, stricter on irreversible side effects
C. Mandatory A2A swarming
D. Semantic memory only

Q20. Max tool-call and timeout limits primarily protect against:

A. Better typography
B. Runaway loops and cost/latency blowups
C. MCP versioning conflicts
D. Vector index rebuilds

Q21. MCP's primary purpose is to:

A. Train new foundation models
B. Standardize how agents/apps connect to tools and context sources
C. Replace databases with prompts
D. Provide only agent-to-agent negotiation

Q22. In MCP terminology, which component exposes tools and resources?

A. MCP Server
B. Vector-only store
C. Browser cookie jar
D. A2A agent card exclusively

Q23. Which statement about adopting MCP is most accurate?

A. Every prototype must start with MCP
B. Ad-hoc local tools can be fine early; MCP value grows with reuse and multi-host integration
C. MCP eliminates authentication needs
D. MCP servers should trust all LLM arguments without validation

Q24. An MCP server exposing arbitrary local file reads to a general employee chatbot creates which immediate concern?

A. Color contrast
B. Over-broad capability and data exfiltration risk without least privilege
C. Lack of multi-agent debate
D. Missing emoji reactions

Q25. MCP is not:

A. A way to expose tools and resources
B. A substitute for an agent-to-agent collaboration protocol

C. Useful for reusable enterprise connectors
D. Something that can sit under an agent host

Q26. A2A primarily addresses:
A. Pixel rendering
B. Agent-to-agent discovery, messaging, and collaboration
C. GPU kernel fusion
D. CSS animations

Q27. Correct complementary use of MCP and A2A is:
A. A2A for database rows; MCP only for agent chat
B. A2A between agents; each agent uses MCP or tools to reach systems of record
C. Never combine them
D. MCP fully replaces orchestrators

Q28. An 'agent card' concept is closest to:
A. A payment card for tokens
B. A public profile of an agent's identity and capabilities
C. A Docker image of PostgreSQL
D. A UI theme pack

Q29. Two agents keep delegating the same task indefinitely. Best mitigation set?
A. Remove schemas and hope
B. Max hop count, clear ownership, structured task contracts, timeouts
C. Increase temperature and parallelism
D. Disable logging

Q30. Fan-out multi-agent pattern means:
A. Only one agent may ever run
B. Multiple agents work (often in parallel) and results are merged
C. Deleting MCP servers
D. Disabling HITL permanently

Q31. Which set is the strongest must-have before production tool-writing agents?
A. Fancy multi-agent debate UI first
B. AuthN/AuthZ, tracing, timeouts, and guardrails on high-risk tools
C. Ten memory types on day one
D. Custom foundation model pretrained from scratch

Q32. A team wants A2A multi-agent planning before tool auth, traces, or evals. What should you recommend?
A. Approve - protocols always come first
B. Reorder: secure tools, observability, and evals first; add multi-agent when the use case forces it
C. Skip tools and only do A2A
D. Replace evals with more prompt adjectives

Q33. Should-have production features typically include:
A. Ignoring failures
B. HITL gates, eval harness, cost controls, and versioning
C. Unauthenticated admin tools
D. No logging by policy

Q34. Choosing low-code visual flows vs code-first graphs should mainly depend on:

A. Which logo looks nicer
B. Team skills, need for complex logic/testability, and enterprise constraints
C. Whether MCP forbids low-code
D. Whether A2A requires Python

Q35. An AI gateway's main job is closest to:

A. Centralizing model access: routing, keys, rate limits, fallbacks
B. Storing salaries in prompts
C. Drawing diagrams only
D. Replacing operational databases

Q36. Which statement correctly relates RAG and agents?

A. RAG alone is always a full agentic system
B. RAG grounds answers with fetched knowledge; agents add planning and tool actions; they often combine
C. Agents cannot use retrieval
D. RAG replaces guardrails

Q37. Observability and evals differ because:

A. They are identical
B. Observability inspects runs; evals systematically score quality and regressions
C. Evals only measure CPU temperature
D. Observability forbids tool-call traces

Q38. Which kill-switch design is most practical for a runaway agent?

A. Hope the model stops
B. Enforce max steps/tool calls, timeouts, budget caps, and admin cancel
C. Delete the company wiki
D. Disable HTTPS

Q39. Channel adapters (Slack/Teams/web) should ideally:

A. Each reimplement business logic separately
B. Reuse the same agent core behind different fronts
C. Bypass auth for convenience
D. Store secrets in chat messages

Q40. You must design an HR assistant for policy Q&A and leave submit with strong audit needs. Strongest architecture statement?

A. One unconstrained autonomous agent with filesystem MCP and no logs
B. Hybrid graph: retrieval for policy; deterministic validated path for leave submit with AuthZ, optional HITL, traces, and golden evals
C. Five A2A agents debating every FAQ
D. Prompt-only with no tools

Q41. A vendor provides an external benefits specialist agent. Your orchestrator should call it for niche questions. Best protocol pairing?

A. Call vendor via A2A; your agent still uses MCP/tools for internal HRIS
B. Put HRIS passwords in the vendor agent card
C. Use only MCP to talk to the vendor agent as if it were a person, with no auth
D. Avoid schemas for cross-org calls

Q42. After a prompt tweak, tool-choice accuracy drops 20% but chat 'sounds nicer.' What was most likely skipped?

A. Font review
B. Evaluation/regression checks on tool selection and task success
C. Immediately adding more agents
D. Removing observability

Q43. Which memory design is most appropriate for a regulated financial agent?

A. Save raw card PANs into semantic memory forever
B. Minimal retention, purpose-limited memory, access control; re-fetch sensitive data from systems of record when needed
C. Mirror production DB into the prompt weekly
D. Disable all session memory so fraud cannot be discussed

Q44. Approving a wire transfer above threshold should be:

A. Fully free-form LLM creativity
B. Deterministic policy and AuthZ (plus HITL), with LLM maybe drafting rationale - not inventing approval
C. Random specialist agent vote without logs
D. MCP resource read of a poem

Q45. A leave bot answers 'You have 14 days left' without any API/tool call. Primary failure mode?

A. Correct tool grounding
B. Hallucination / ungrounded answer instead of live system data
C. Perfect AuthZ
D. Deterministic HRIS write

Q46. For structured tool-argument generation, which setting is usually safer?

A. Very high temperature for creativity
B. Lower temperature for more stable structured outputs
C. No schema, only free prose
D. Disable all validation

Q47. Soft policy text in the system prompt is insufficient for 'cannot approve own expenses' because:

A. Prompts are illegal
B. Models can ignore or be manipulated; enforcement belongs in code and tool AuthZ
C. Expenses never need rules
D. Only A2A can enforce policy

Q48. Submitting leave twice creates two records. Which concept was missing?

A. Temperature tuning
B. Idempotency and/or duplicate detection for write operations
C. Vector embeddings
D. Agent cards

Q49. Why unit-test tools without calling the LLM?

A. LLMs are always free
B. Tools are deterministic business logic; tests catch contract bugs cheaply and stably
C. Unit tests replace AuthZ
D. Frameworks forbid LLM tests forever

Q50. Which interview-ready summary best reflects mature enterprise practice?

A. Always maximize agent count and autonomy on day one
B. Build hybrid systems: rails for compliance-critical steps, LLM agents for flexible reasoning, narrow authenticated tools,

guardrails/HITL on writes, traces plus golden evals; multi-agent/A2A only when specialist ownership requires it

C. Put all security rules only in prompts and skip evals

D. Prefer one mega-tool with admin rights for speed

Answer Key

Review only after attempting the full set. Each item lists the correct option.

Q1:B Q2:B Q3:B Q4:C Q5:B
Q6:B Q7:B Q8:B Q9:B Q10:B
Q11:B Q12:B Q13:B Q14:B Q15:B
Q16:B Q17:B Q18:B Q19:B Q20:B
Q21:B Q22:A Q23:B Q24:B Q25:B
Q26:B Q27:B Q28:B Q29:B Q30:B
Q31:B Q32:B Q33:B Q34:B Q35:A
Q36:B Q37:B Q38:B Q39:B Q40:B
Q41:A Q42:B Q43:B Q44:B Q45:B
Q46:B Q47:B Q48:B Q49:B Q50:B

Brief rationale themes (for discussion)

- Prefer architectural specificity over buzzwords; agentic is a spectrum.
- Hybrid orchestration is the usual enterprise fit for compliance + flexibility.
- Enforce AuthZ and irreversible actions in the tool layer, not prompts alone.
- MCP connects agents to tools/context; A2A connects agents to agents.
- Ship tools, guardrails, traces, and evals before exotic multi-agent designs.
- Ground live facts with tools; use evals to catch regressions after prompt changes.

Source basis: I2I Agentic AI Starter Guide (FDE) study pack. For internal learning use.